



Deep Learning

AAKASH GHOSH
19MS129



Presentation made for the completion of Course [MA5115](#)

Traditional ML



Traditional Machine learning

The older/traditional way of applying machine learning consisted of the following steps:

- **Feature Extraction:** Features which can be used to discriminate between classes is captured. This step usually requires in-field knowledge about the problem.

Traditional Machine learning

The older/traditional way of applying machine learning consisted of the following steps:

- **Feature Extraction:** Features which can be used to discriminate between classes is captured. This step usually requires in-field knowledge about the problem.
- **Model Selection:** A model is selected which trains on the extracted features. Ensemble methods can be used to boost performance

Traditional Machine learning

The older/traditional way of applying machine learning consisted of the following steps:

- **Feature Extraction:** Features which can be used to discriminate between classes is captured. This step usually requires in-field knowledge about the problem.
- **Model Selection:** A model is selected which trains on the extracted features. Ensemble methods can be used to boost performance
- **Cross-validation/Tested:** The model is tested/cross-validated on with held data

Problems with the traditional approach

- **Feature Extraction:** This step requires in-field knowledge. It is very difficult to study a whole new branch of knowledge for a single problem.

Problems with the traditional approach

- **Feature Extraction:** This step requires in-field knowledge. It is very difficult to study a whole new branch of knowledge for a single problem.
- **Amount of Data:** In the current era, the amount of data sometimes is simply so large that meaningful features are hard to extract

Problems with the traditional approach

- **Feature Extraction:** This step requires in-field knowledge. It is very difficult to study a whole new branch of knowledge for a single problem.
- **Amount of Data:** In the current era, the amount of data sometimes is simply so large that meaningful features are hard to extract
- **Unorganized Data:** Feature extraction is hard in unorganized data (such as a literary works).

The idea behind deep learning

OH, HEY, YOU ORGANIZED
OUR PHOTO ARCHIVE!

YEAH, I TRAINED A NEURAL
NET TO SORT THE UNLABELED
PHOTOS INTO CATEGORIES.

WHOA! NICE WORK!



ENGINEERING TIP:
WHEN YOU DO A TASK BY HAND,
YOU CAN TECHNICALLY SAY YOU
TRAINED A NEURAL NET TO DO IT.

Ideas behind a neural network

- **Pattern recognition:** Almost all the problems with traditional ML lies in proper feature selection. We now take a different approach: We train the machine to learn what the necessary patterns/features are.

Ideas behind a neural network

- **Pattern recognition:** Almost all the problems with traditional ML lies in proper feature selection. We now take a different approach: We train the machine to learn what the necessary patterns/features are.
- **Inspiration:** For this task we take inspiration from the best pattern recognizing machine that we know:


Ideas behind a neural network

- **Pattern recognition:** Almost all the problems with traditional ML lies in proper feature selection. We now take a different approach: We train the machine to learn what the necessary patterns/features are.
- **Inspiration:** For this task we take inspiration from the best pattern recognizing machine that we know: **our brains**.

Ideas behind a neural network

- **Pattern recognition:** Almost all the problems with traditional ML lies in proper feature selection. We now take a different approach: We train the machine to learn what the necessary patterns/features are.
- **Inspiration:** For this task we take inspiration from the best pattern recognizing machine that we know: **our brains**.
- **Overview:** We make a perceptron, which is essentially the digital analog of a neuron. We connect multiple neurons together(similar to our brain) and look at what pattern and amount of activation leads to best result.

A practical comparison with traditional ML




Comparison based on amount of data features.

Src: Neural Networks and Deep Learning: A Textbook, Charu C.

Aggarwal

- Traditional ML models show better prediction when the amount of features involved is small. Features can be individually engineered and interpreted.

A practical comparison with traditional ML




Comparison based on amount of data features.

Src: Neural Networks and Deep Learning: A Textbook, Charu C.


Aggarwal

- Traditional ML models show better prediction when the amount of features involved is small. Features can be individually engineered and interpreted.
- Deep learning models are better when data is unstructured or there are a lot of features which need to be considered. With proper construction and training almost any decision boundaries can be learned.

The Perceptron



Non-Linearity is the key




A perceptron

Src: MIT Introduction to Deep Learning, 6.S191, Lec-1

- The idea behind a perceptron is that a perceptron will give a high value as output when it recognises a certain feature.

Non-Linearity is the key




A perceptron

Src: MIT Introduction to Deep Learning, 6.S191, Lec-1

- The idea behind a perceptron is that a perceptron will give a high value as output when it recognises a certain feature.
- Mathematically for an input vector X it outputs $\hat{Y} = \sigma(W^T X + b_0)$

Non-Linearity is the key




A perceptron

Src: MIT Introduction to Deep Learning, 6.S191, Lec-1

- The idea behind a perceptron is that a perceptron will give a high value as output when it recognises a certain feature.
- Mathematically for an input vector X it outputs $\hat{Y} = \sigma(W^T X + b_0)$
- Multiple perceptrons are wired together (again, much like our brain) to identify richer features.

Non-Linearity is the key




A perceptron

Src: MIT Introduction to Deep Learning, 6.S191, Lec-1

- The idea behind a perceptron is that a perceptron will give a high value as output when it recognises a certain feature.
- Mathematically for an input vector X it outputs $\hat{Y} = \sigma(W^T X + b_0)$
- Multiple perceptrons are wired together (again, much like our brain) to identify richer features.
- At the very core, Machine Learning works by constructing proper decision boundaries.

Non-Linearity is the key




A perceptron


Src: MIT Introduction to Deep Learning, 6.S191, Lec-1

- The idea behind a perceptron is that a perceptron will give a high value as output when it recognises a certain feature.
- Mathematically for an input vector X it outputs $\hat{Y} = \sigma(W^T X + b_0)$
- Multiple perceptrons are wired together (again, much like our brain) to identify richer features.
- At the very core, Machine Learning works by constructing proper decision boundaries.
- A very simple calculation shows if σ is linear then the decision boundary is also linear. It therefore makes sense to use non-linear σ .

Non-Linearity is the key





- As shown in the figure on left, classes which can't be separated by linear boundaries can be separated by non-linear functions. (Think SVM but the kernels are learned.)



Non-linearity in action

Src: *Neural Networks and Deep Learning: A Textbook*, Charu C. Aggarwal

Non-Linearity is the key




Non-linearity in action

Src: Neural Networks and Deep Learning: A Textbook, Charu C. Aggarwal

- As shown in the figure on left, classes which can't be separated by linear boundaries can be separated by non-linear functions.(Think SVM but the kernels are learned.)
- It can be theoretically shown that almost all boundary function can be separated by a 2 layered neural network.

Increasing depth




Why is depth needed

Src: Neural Networks and Deep Learning: A Textbook, Charu C.

Aggarwal

- While a single hidden layer is enough, making a deep neural network allows us to have more complex decision boundaries with relatively less number of nodes

Increasing depth



Why is depth needed

Src: Neural Networks and Deep Learning: A Textbook, Charu C.

Aggarwal


- While a single hidden layer is enough, making a deep neural network allows us to have more complex decision boundaries with relatively less number of nodes
- It should be noted how non-linearity discussed above plays an important role: Irrespective of number of layers, composition of linear functions is linear. On the other hand compositions of non-linear function can lead to richer families of functions. (Ref: <https://www.desmos.com/calculator/m645ggyl2i>)

Training a neural network

Loss functions, Backpropagation, Reversemode Auto-Diff, Practical problems and hardware considerations




Loss Function



*A slice of the loss landscape of resnet-110 with no skip connections
Src: Visualizing the Loss Landscape of Neural Nets, Hao Li, Zheng
Xu, Gavin Taylor, Christoph Studer, Tom Goldstein*

- We wish to calculate weights $w_i, 1 \leq i \leq n$ so that the predicted values \hat{Y} are as close as possible (to an extent) to the actual values Y .


Loss Function



*A slice of the loss landscape of resnet-110 with no skip connections
Src: Visualizing the Loss Landscape of Neural Nets, Hao Li, Zheng
Xu, Gavin Taylor, Christoph Studer, Tom Goldstein*

- We wish to calculate weights $w_i, 1 \leq i \leq n$ so that the predicted values \hat{Y} are as close as possible (to an extent) to the actual values Y . To do this we try to minimize a loss function $\mathcal{L}(\{\hat{Y}_i\}, \{Y_i\})$


Loss Function



*A slice of the loss landscape of resnet-110 with no skip connections
Src: Visualizing the Loss Landscape of Neural Nets, Hao Li, Zheng
Xu, Gavin Taylor, Christoph Studer, Tom Goldstein*

- We wish to calculate weights $w_i, 1 \leq i \leq n$ so that the predicted values \hat{Y} are as close as possible (to an extent) to the actual values Y . To do this we try to minimize a loss function $\mathcal{L}(\{\hat{Y}_i\}, \{Y_i\})$
- Common choices of \mathcal{L} include MSE for continuous variables and cross-entropy for categorical variables.

Gradient descent



Gradient descent


Src: <https://towardsdatascience.com/an-intuitive-explanation-of-gradient-descent-83adf68c9c33>

- To find the correct set of weights, we use a greedy approach. We check the surrounding landscape of the weight (i.e. calculate the gradient) and take a step in the direction which leads to maximum decrease in \mathcal{L} .

Mathematically, we have:

$$w_i = w_i - \eta \frac{\partial \mathcal{L}}{\partial w_i}$$

Gradient descent



Gradient descent

Src: <https://towardsdatascience.com/an-intuitive-explanation-of-gradient-descent-83adf68c9c33>

- To find the correct set of weights, we use a greedy approach. We check the surrounding landscape of the weight (i.e. calculate the gradient) and take a step in the direction which leads to maximum decrease in \mathcal{L} .


Mathematically, we have:

$$w_i = w_i - \eta \frac{\partial \mathcal{L}}{\partial w_i}$$

- η is known as the learning rate.

Learning Rate


- Fixing η is quite tricky.



Gradient descent

Src: <https://towardsdatascience.com/an-intuitive-explanation-of-gradient-descent-83adf68c9c33>

Learning Rate




Gradient descent

Src: <https://towardsdatascience.com/an-intuitive-explanation-of-gradient-descent-83adf68c9c33>

- Fixing η is quite tricky.
- If η is small we get stuck at local minima

Learning Rate



Gradient descent

Src: <https://towardsdatascience.com/an-intuitive-explanation-of-gradient-descent-83adf68c9c33>

- Fixing η is quite tricky.
- If η is small we get stuck at local minima
- If η is large we overshoot our targets and never converge.
- The best way to do things is to use a adaptive learn rate. Some methods(parametric, non-parametric and hybrid are discussed later, once we cover backpropagation)

Backpropagation

Consider an acyclic directed graph.

- We want to calculate the partial derivative of a node with respect to the other.

Backpropagation

Consider an acyclic directed graph.

- We want to calculate the partial derivative of a node with respect to the other.
- The way to do this is to use the chain rule

Backpropagation

Consider an acyclic directed graph.

- We want to calculate the partial derivative of a node with respect to the other.
- The way to do this is to use the chain rule
- When we use the chain rule, we sum the partial derivative along all the paths joining the two nodes
- It is easy to see that the computation needed explodes as the depth increases
- The way to navigate this problem is to use dynamic programming.

References