



Deep Learning

AAKASH GHOSH
19MS129



Presentation made for the completion of Course [MA5115](#)

Traditional ML



Traditional Machine learning

The older/traditional way of applying machine learning consisted of the following steps:

- **Feature Extraction:** Features which can be used to discriminate between classes is captured. This step usually requires in-field knowledge about the problem.

Traditional Machine learning

The older/traditional way of applying machine learning consisted of the following steps:

- **Feature Extraction:** Features which can be used to discriminate between classes is captured. This step usually requires in-field knowledge about the problem.
- **Model Selection:** A model is selected which trains on the extracted features. Ensemble methods can be used to boost performance

Traditional Machine learning

The older/traditional way of applying machine learning consisted of the following steps:

- **Feature Extraction:** Features which can be used to discriminate between classes is captured. This step usually requires in-field knowledge about the problem.
- **Model Selection:** A model is selected which trains on the extracted features. Ensemble methods can be used to boost performance
- **Cross-validation/Tested:** The model is tested/cross-validated on with held data

Problems with the traditional approach

- **Feature Extraction:** This step requires in-field knowledge. It is very difficult to study a whole new branch of knowledge for a single problem.

Problems with the traditional approach

- **Feature Extraction:** This step requires in-field knowledge. It is very difficult to study a whole new branch of knowledge for a single problem.
- **Amount of Data:** In the current era, the amount of data sometimes is simply so large that meaningful features are hard to extract

Problems with the traditional approach

- **Feature Extraction:** This step requires in-field knowledge. It is very difficult to study a whole new branch of knowledge for a single problem.
- **Amount of Data:** In the current era, the amount of data sometimes is simply so large that meaningful features are hard to extract
- **Unorganized Data:** Feature extraction is hard in unorganized data (such as a literary works).

The idea behind deep learning

OH, HEY, YOU ORGANIZED
OUR PHOTO ARCHIVE!

YEAH, I TRAINED A NEURAL
NET TO SORT THE UNLABELED
PHOTOS INTO CATEGORIES.

WHOA! NICE WORK!



ENGINEERING TIP:
WHEN YOU DO A TASK BY HAND,
YOU CAN TECHNICALLY SAY YOU
TRAINED A NEURAL NET TO DO IT.

Ideas behind a neural network

- **Pattern recognition:** Almost all the problems with traditional ML lies in proper feature selection. We now take a different approach: We train the machine to learn what the necessary patterns/features are.

Ideas behind a neural network

- **Pattern recognition:** Almost all the problems with traditional ML lies in proper feature selection. We now take a different approach: We train the machine to learn what the necessary patterns/features are.
- **Inspiration:** For this task we take inspiration from the best pattern recognizing machine that we know:


Ideas behind a neural network

- **Pattern recognition:** Almost all the problems with traditional ML lies in proper feature selection. We now take a different approach: We train the machine to learn what the necessary patterns/features are.
- **Inspiration:** For this task we take inspiration from the best pattern recognizing machine that we know: **our brains**.

Ideas behind a neural network

- **Pattern recognition:** Almost all the problems with traditional ML lies in proper feature selection. We now take a different approach: We train the machine to learn what the necessary patterns/features are.
- **Inspiration:** For this task we take inspiration from the best pattern recognizing machine that we know: **our brains**.
- **Overview:** We make a perceptron, which is essentially the digital analog of a neuron. We connect multiple neurons together(similar to our brain) and look at what pattern and amount of activation leads to best result.

A practical comparison with traditional ML




- Traditional ML models show better prediction when the amount of features involved is small. Features can be individually engineered and interpreted.

Comparison based on amount of data features.

Src: Neural Networks and Deep Learning: A Textbook, Charu C. Aggarwal

A practical comparison with traditional ML




Comparison based on amount of data features.


Src: Neural Networks and Deep Learning: A Textbook, Charu C. Aggarwal

- Traditional ML models show better prediction when the amount of features involved is small. Features can be individually engineered and interpreted.
- Deep learning models are better when data is unstructured or there are a lot of features which need to be considered. With proper construction and training almost any decision boundaries can be learned.

The Perceptron



Non-Linearity is the key




A perceptron

Src: MIT Introduction to Deep Learning, 6.S191, Lec-1

- The idea behind a perceptron is that a perceptron will give a high value as output when it recognises a certain feature.

Non-Linearity is the key




A perceptron

Src: MIT Introduction to Deep Learning, 6.S191, Lec-1

- The idea behind a perceptron is that a perceptron will give a high value as output when it recognises a certain feature.
- Mathematically for an input vector X it outputs $\hat{Y} = \sigma(W^T X + b_0)$

Non-Linearity is the key




A perceptron

Src: MIT Introduction to Deep Learning, 6.S191, Lec-1

- The idea behind a perceptron is that a perceptron will give a high value as output when it recognises a certain feature.
- Mathematically for an input vector X it outputs $\hat{Y} = \sigma(W^T X + b_0)$
- Multiple perceptrons are wired together (again, much like our brain) to identify richer features.

Non-Linearity is the key




A perceptron

Src: MIT Introduction to Deep Learning, 6.S191, Lec-1

- The idea behind a perceptron is that a perceptron will give a high value as output when it recognises a certain feature.
- Mathematically for an input vector X it outputs $\hat{Y} = \sigma(W^T X + b_0)$
- Multiple perceptrons are wired together (again, much like our brain) to identify richer features.
- At the very core, Machine Learning works by constructing proper decision boundaries.

Non-Linearity is the key




A perceptron


Src: MIT Introduction to Deep Learning, 6.S191, Lec-1

- The idea behind a perceptron is that a perceptron will give a high value as output when it recognises a certain feature.
- Mathematically for an input vector X it outputs $\hat{Y} = \sigma(W^T X + b_0)$
- Multiple perceptrons are wired together (again, much like our brain) to identify richer features.
- At the very core, Machine Learning works by constructing proper decision boundaries.
- A very simple calculation shows if σ is linear then the decision boundary is also linear. It therefore makes sense to use non-linear σ .

Non-Linearity is the key





- As shown in the figure on left, classes which can't be separated by linear boundaries can be separated by non-linear functions. (Think SVM but the kernels are learned.)



Non-linearity in action

Src: *Neural Networks and Deep Learning: A Textbook*, Charu C. Aggarwal

Non-Linearity is the key




Non-linearity in action

Src: *Neural Networks and Deep Learning: A Textbook*, Charu C. Aggarwal

- As shown in the figure on left, classes which can't be separated by linear boundaries can be separated by non-linear functions.(Think SVM but the kernels are learned.)
- It can be theoretically shown that almost all boundary function can be separated by a 2 layered neural network.

Increasing depth




Why is depth needed

Src: Neural Networks and Deep Learning: A Textbook, Charu C. Aggarwal

- While a single hidden layer is enough, making a deep neural network allows us to have more complex decision boundaries with relatively less number of nodes

Increasing depth



Why is depth needed

Src: *Neural Networks and Deep Learning: A Textbook*, Charu C. Aggarwal


- While a single hidden layer is enough, making a deep neural network allows us to have more complex decision boundaries with relatively less number of nodes
- It should be noted how non-linearity discussed above plays an important role: Irrespective of number of layers, composition of linear functions is linear. On the other hand compositions of non-linear function can lead to richer families of functions. (Ref: <https://www.desmos.com/calculator/m645ggyl2i>)

Training a neural network

Loss functions, Backpropagation, Reversemode Auto-Diff, Practical problems and hardware considerations



Loss Function




A slice of the loss landscape of resnet-110 with no skip connections

Src: Visualizing the Loss Landscape of Neural Nets, Hao Li, Zheng Xu, Gavin Taylor, Christoph Studer, Tom Goldstein

- We wish to calculate weights $w_i, 1 \leq i \leq n$ so that the predicted values \hat{Y} are as close as possible (to an extent) to the actual values Y .

Loss Function




A slice of the loss landscape of resnet-110 with no skip connections

Src: Visualizing the Loss Landscape of Neural Nets, Hao Li, Zheng Xu, Gavin Taylor, Christoph Studer, Tom Goldstein

- We wish to calculate weights $w_i, 1 \leq i \leq n$ so that the predicted values \hat{Y} are as close as possible (to an extent) to the actual values Y . To do this we try to minimize a loss function $\mathcal{L}(\{\hat{Y}_i\}, \{Y_i\})$

Loss Function




A slice of the loss landscape of resnet-110 with no skip connections

Src: Visualizing the Loss Landscape of Neural Nets, Hao Li, Zheng Xu, Gavin Taylor, Christoph Studer, Tom Goldstein

- We wish to calculate weights $w_i, 1 \leq i \leq n$ so that the predicted values \hat{Y} are as close as possible (to an extent) to the actual values Y . To do this we try to minimize a loss function $\mathcal{L}(\{\hat{Y}_i\}, \{Y_i\})$
- Common choices of \mathcal{L} include MSE for continuous variables and cross-entropy for categorical variables.

Gradient descent




Gradient descent

Src: <https://towardsdatascience.com/an-intuitive-explanation-of-gradient-descent-83adf68c9c33>

- To find the correct set of weights, we use a greedy approach. We check the surrounding landscape of the weight (i.e. calculate the gradient) and take a step in the direction which leads to maximum decrease in \mathcal{L} . Mathematically, we have:

$$w_i = w_i - \eta \frac{\partial \mathcal{L}}{\partial w_i}$$

Gradient descent



Gradient descent


Src: <https://towardsdatascience.com/an-intuitive-explanation-of-gradient-descent-83adf68c9c33>

- To find the correct set of weights, we use a greedy approach. We check the surrounding landscape of the weight (i.e. calculate the gradient) and take a step in the direction which leads to maximum decrease in \mathcal{L} . Mathematically, we have:

$$w_i = w_i - \eta \frac{\partial \mathcal{L}}{\partial w_i}$$

- η is known as the learning rate.

Learning Rate




- Fixing η is quite tricky.

Variations in learning rates

*Src: Hands-On Machine Learning with Scikit-Learn, Keras,
and TensorFlow Concepts, Tools, and Techniques to Build
Intelligent Systems-O'Reilly Media, Inc, Aurélien Géron*

Learning Rate




Variations in learning rates

Src: *Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow Concepts, Tools, and Techniques to Build Intelligent Systems*-O'Reilly Media, Inc, Aurélien Géron

- Fixing η is quite tricky.
- If η is small we get stuck at local minima

Learning Rate



Variations in learning rates

Src: Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow Concepts, Tools, and Techniques to Build Intelligent Systems-O'Reilly Media, Inc, Aurélien Géron

- Fixing η is quite tricky.
- If η is small we get stuck at local minima
- If η is large we overshoot our targets and never converge.
- The best way to do things is to use an adaptive learn rate. Some methods (parametric, non-parametric and hybrid) are discussed later, once we cover backpropagation)

Backpropagation

Consider a neural network to be an acyclic directed graph

- We want to calculate the partial derivative of a node with respect to the other.

Backpropagation


Consider a neural network to be an acyclic directed graph

- We want to calculate the partial derivative of a node with respect to the other.
- The way to do this is to use the chain rule

Backpropagation

Consider a neural network to be an acyclic directed graph

- We want to calculate the partial derivative of a node with respect to the other.
- The way to do this is to use the chain rule
- When we use the chain rule, we sum the partial derivative along all the paths joining the two nodes
- It is easy to see that the computation needed explodes as the depth increases
- The way to navigate this problem is to use dynamic programming.



$$\frac{\partial h}{\partial x} = \frac{\partial h}{\partial f} \cdot \frac{\partial f}{\partial x} + \frac{\partial h}{\partial g} \cdot \frac{\partial g}{\partial x}$$

Backpropagation

Roughly, the algorithm works as follows:


- Initialize the weights

Backpropagation

Roughly, the algorithm works as follows:

- Initialize the weights
- Calculate \mathcal{L} . Note, that this step occurs from the input towards the output and is known as the forward phase.
- Calculate gradient. This process occurs from the output towards the input and is known as the backward phase.


Reverse Mode Auto Differentiation



$$\begin{aligned}\frac{\partial f}{\partial x} &= \frac{\partial f}{\partial x_3} \cdot \frac{\partial x_3}{\partial x} + \frac{\partial f}{\partial y_3} \cdot \frac{\partial y_3}{\partial x} \\ &= \frac{\partial f}{\partial x_3} \cdot \left(\frac{\partial x_3}{\partial x_2} \frac{\partial x_2}{\partial x} + \frac{\partial x_3}{\partial y_2} \frac{\partial y_2}{\partial x} \right) + \frac{\partial f}{\partial y_3} \cdot \left(\frac{\partial y_3}{\partial x_2} \frac{\partial x_2}{\partial x} + \frac{\partial y_3}{\partial y_2} \frac{\partial y_2}{\partial x} \right)\end{aligned}$$


The whole idea relies on the fact that while calculating gradients, parts of the calculation are repeated.

Reverse Mode Auto Differentiation




Sum is taken over the product along each path, over all the paths. It is trivial to note that parts of the paths are repeated.

Reverse Mode Auto Differentiation



Sum is taken over the product along each path, over all the paths. It is trivial to note that parts of the paths are repeated. In the toy example on the left, the paths are essentially the same, varying only for the second to last node.

Stochastic Gradient Descent




Gradient descent

Src: <https://towardsdatascience.com/an-intuitive-explanation-of-gradient-descent-83adf68c9c33>

- Fixing η is quite tricky.

Stochastic Gradient Descent




Gradient descent

Src: <https://towardsdatascience.com/an-intuitive-explanation-of-gradient-descent-83adf68c9c33>

- Fixing η is quite tricky.
- If η is small we get stuck at local minima

Stochastic Gradient Descent



Gradient descent

Src: <https://towardsdatascience.com/an-intuitive-explanation-of-gradient-descent-83adf68c9c33>

- Fixing η is quite tricky.
- If η is small we get stuck at local minima
- If η is large we overshoot our targets and never converge.
- The best way to do things is to use an adaptive learn rate. Some methods (parametric, non-parametric and hybrid) are discussed later, once we cover backpropagation)

References